

— Session 09: NLP basics

wifi: GA-Guest, yellowpencil

```
cd ~/Documents/ga-ldn-ds37  
git commit -am "your commit message here"  
git pull
```



Today's session plan

1800-1810	Standup
1830-1900	Linear algebra review
1900-1920	Break
1920-2100	Project troubleshooting

At the end of the session, you will be able to ...

Vectorise text using count vectorisers and TF-IDF

Use pre-processing methods including stemming and removing stopwords

Implement text classification with Naive Bayes

Develop your independent project work further

Data Science Part Time

What is NLP?

Natural language processing

Natural language processing is a field concerned with:

- Using computers to process human language and text
- Making sense of human knowledge stored as unstructured text

Natural language processing has many applications in data science, including:

- Text classification
- Topic modelling or clustering
- Translation
- Text summarisation or simplification
- Sentiment analysis
- Speech to text
- Text to speech

Natural language processing

Many natural language processing tasks will involve the following pre-processing techniques:

Tokenization Breaking a sentence into words

Stopword removal Removing and/it/they etc and other grammatically non-useful words

Stemming and lemmatization Reducing words back to their original roots

Part of speech tagging Labelling nouns, verbs, adjectives

Vectorisation Representing text data as a numerical vector

Why is natural language processing hard?

Here are just a few of the reasons why NLP is challenging:

Ambiguity

Non standard English

Tricky entity names

Newly coined words

Idioms



Computers Out: Getting started with NLP



Now let's open ds37-09-01.ipynb to get started

Some Terminology

Document

A single observation in a dataset. This could be a single review in a database of customer reviews, a single tweet in a set of tweets, etc

Corpus

A collection of documents

Token

A single word in a document. Sometimes a token can be a single phrase, of a specified length. For example, we might want to split a document into two-word tokens.

Stopwords

Stop words are some of the most common words in a language. They are used so that a sentence makes sense grammatically, such as prepositions and determiners, e.g., "to," "the," "and." However, they are so commonly used that they are generally worthless for predicting the class of a document.

```
> stopwords("english")
[1] "i"      "me"      "my"      "myself"  "we"
[6] "our"    "ours"    "ourselves" "you"     "your"
[11] "yours"  "yourself" "yourselves" "he"      "him"
[16] "his"    "himself" "she"      "her"     "hers"
[21] "herself" "it"      "its"      "itself"  "they"
[26] "them"   "their"   "theirs"   "themselves" "what"
[31] "which"  "who"     "whom"    "this"    "that"
[36] "these"  "those"   "am"      "is"      "are"
[41] "was"    "were"    "be"      "been"    "being"
[46] "have"   "has"     "had"     "having"  "do"
```

Stemming and lemmatization

These are both ways of intelligently reducing the number of features by grouping together (hopefully) related words.

Stemming is the process of removing common endings from sentences, such as "s", "es", "ly", "ing", and "ed". This reduce a word to its base/stem/root form.

Lemmatization is a more refined process that uses specific language and grammar rules to derive the root of a word. This is useful for words that do not share an obvious root such as "better" and "best".

Lemmatization	Stemming
shouted → shout	badly → bad
best → good	computing → comput
better → good	computed → comput
good → good	wipes → wip
wiping → wipe	wiped → wip
hidden → hide	wiping → wip

Bag of words model

Converts a corpus of text to a **term-document matrix**, where every row corresponds to a **document** and every column is a feature or **unique word**.

The value of each element in the matrix is either a binary indicator, marking the presence of that word in the document, or a word count.

How can we convert a corpus of text to a matrix?

	it	is	puppy	cat	pen	a	this
it is a puppy	1	1	1	0	0	1	0
it is a kitten	1	1	0	0	0	1	0
it is a cat	1	1	0	1	0	1	0
that is a dog and this is a pen	0	2	0	0	1	2	1
it is a matrix	1	1	0	0	0	1	0

Vectorising

There are a few different ways of vectorising text documents. One simple method is **count vectorising**. Each matrix entry is the count of a particular word in a particular document.

	it	is	puppy	cat	pen	a	this
it is a puppy	1	1	1	0	0	1	0
it is a kitten	1	1	0	0	0	1	0
it is a cat	1	1	0	1	0	1	0
that is a dog and this is a pen	0	2	0	0	1	2	1
it is a matrix	1	1	0	0	0	1	0



Solo Exercise:

Count vectorising



Take the following sentences and go through the steps below, by hand:

The cat sat on the mat

The rat was sat on the cat

The mat is red

The cat is black

1. Compile the **vocabulary** for this corpus. That's the set of unique words across the whole corpus, with no repeat words.
2. On paper, compile the term-document matrix for this corpus.

Naive Bayes

Once we've vectorised our text, we use Bayes' theorem to learn the relationship between a document's features (i.e. the words in it) and its label.

Bayes' Theorem, is a way to assess probabilities, using prior information to make more accurate predictions.

$$P(A | B) = \frac{P(B | A) \times P(A)}{P(B)}$$

- $P(A | B)$: Probability of Event A occurring given Event B has occurred.
- $P(B | A)$: Probability of Event B occurring given Event A has occurred.
- $P(A)$: Probability of Event A occurring.
- $P(B)$: Probability of Event B occurring.



Solo Exercise:

Bayes' theorem



At a GP surgery, 10% of patients are prescribed painkillers.

Overall, five percent of the surgery's patients are addicted to painkillers.

Out of all the people prescribed pain pills, 8% are addicts.

If a patient is an addict, what is the probability that they will be prescribed pain pills?

Bayes' theorem

Step 1 Figure out what your event “A” is from the question. That information is in the italicized part of this particular question. The event that happens first (A) is being prescribed pain pills. That’s given as 10%.

Step 2 Figure out what your event “B” is from the question. That information is also in the italicized part of this particular question. Event B is being an addict. That’s given as 5%.

Step 3 Figure out what the probability of event B (Step 2) given event A (Step 1). In other words, find what (B|A) is. We want to know “Given that people are prescribed pain pills, what’s the probability they are an addict?” That is given in the question as 8%, or .8.

Step 4 Insert your answers from Steps 1, 2 and 3 into the formula and solve.

$$P(A|B) = P(B|A) * P(A) / P(B) = (0.08 * 0.1) / 0.05 = 0.16$$

The probability of an addict being prescribed pain pills is 0.16 (16%).

Bayes' theorem

How does Bayes' theorem apply to our text classification problem? Imagine we're trying to classify emails into two classes: 'spam' or 'ham'. Our first email contains just three words: 'send money now' so the probability that it's spam is as follows

$$P(\text{spam} \mid \text{send money now}) = \frac{P(\text{send money now} \mid \text{spam}) \times P(\text{spam})}{P(\text{send money now})}$$

We assume our features are **conditionally independent** so our calculation simplifies to:

We know what each of these terms is from the training data

$$P(\text{spam} \mid \text{send money now}) \approx \frac{P(\text{send} \mid \text{spam}) \times P(\text{money} \mid \text{spam}) \times P(\text{now} \mid \text{spam}) \times P(\text{spam})}{P(\text{send money now})}$$

Bayes' Theorem

Imagine we plug these numbers in from our training data:

$$P(\textit{spam} \mid \text{send money now}) \approx \frac{0.2 \times 0.1 \times 0.1 \times 0.9}{P(\text{send money now})} = \frac{0.0018}{P(\text{send money now})}$$

Now we repeat the calculation for our **other** class, ham:

$$P(\textit{ham} \mid \text{send money now}) \approx \frac{0.05 \times 0.01 \times 0.1 \times 0.1}{P(\text{send money now})} = \frac{0.000005}{P(\text{send money now})}$$

All we care about is which class has a higher probability, so we can ignore the denominator and pick the highest probability; in this case that's **spam**.

Bayes' Theorem

The key takeaways from Bayes' are:

- The "naive" assumption of Naive Bayes (that the features are conditionally independent) is critical to making these calculations simple.
- The normalization constant (the denominator) can be ignored since it's the same for all classes.

Advantages of Naive Bayes

Model training and prediction are very fast.

It's somewhat interpretable.

No fine tuning is required.

Disadvantages of Naive Bayes

If "spam" is dependent on non-independent combinations of individual words, it may not work well.

Correlated features can be problematic (due to the independence assumption).

Conditional independence

The events R and B are conditionally independent [given Y] if and only if, given knowledge of whether Y occurs, knowledge of whether R occurs provides no information on the likelihood of B occurring, and knowledge of whether B occurs provides no information on the likelihood of R occurring.

For example

Let the two events be the probabilities of persons A and B getting home in time for dinner, and the third event is the fact that a snow storm hit the city. While both A and B have a lower probability of getting home in time for dinner, the lower probabilities will still be independent of each other. That is, the knowledge that A is late does not tell you whether B will be late. (They may be living in different neighborhoods, traveling different distances, and using different modes of transportation.)

TF-IDF

There are alternative methods for vectorising text; one of these is TF-IDF.

TF-IDF analyses the uniqueness of words between documents to find distinguishing characteristics.

Term frequency–inverse document frequency (TF–IDF) computes the "relative frequency" with which a word appears in a document, compared to its frequency across all documents.

It's more useful than "term frequency" for identifying "important" words in each document (high frequency in that document, low frequency in other documents). The TF-IDF score for a given word in a given document is:

$$\frac{\text{number of times the word is in this document}}{\text{number of documents the word is in}}$$



Solo Exercise: TF-IDF



Given the following documents, compile a term-document matrix using TF-IDF. **Before** you start, which word do you think will have the highest TF-IDF score? Why?

The cat and dog sat

The dog and cat sat

The cat sat and sat

The cat killed the dog

Intro to Python



Let's Review

At the end of the session, you will be able to ...

Vectorise text using count vectorisers and TF-IDF

Use pre-processing methods including stemming and removing stopwords

Implement text classification with Naive Bayes

Develop your independent project work further

Coming up next time...

- Linear regression



